

Vijay Chevireddi

College Park, MD • vijaydevreddychevireddi@gmail.com • 240-940-9294 • [GitHub](#) • [LinkedIn](#) • [Portfolio](#)

EDUCATION

University of Maryland

Master of Engineering (M.Eng.) – Robotics / GPA – 3.77

Aug 2023 – May 2025

College Park, MD

Osmania University

Bachelor of Engineering – Mechanical Engineering / GPA – 9.05/10

Aug 2018 – June 2022

Hyderabad, India

COURSEWORK

Multimodal Foundation Models (CNN, Transformers, RoPE, Flash Attention, Fine Tuning, MAMBA, GPT-1,2,3, Agentic AI, Vision Lang. Models, Variational Auto Encoder, Diffusion Models, Flow Matching, SWIN Transformer, RAG, CLIP, DINO, LLaVA), Natural Language Processing (Audio Signal Processing for ML, RLHF, PPO, DPO, AI Alignment, AI Distillation, BERT, ELMO, RNN, LSTM), Robot Learning (Reinforcement Learning), Perception for Autonomous Robots (Computer Vision), Planning for Autonomous Robots, Control Systems

SKILLS

- **Programming Languages:** Python, C/C++, SQL, MATLAB
- **Tools & Frameworks:** PyTorch, TensorFlow, AWS, SageMaker, SQL, OpenAI Gym, MuJoCo, Langchain, Langgraph, Hugging Face, Docker, Scikit-learn, xgboost, NLTK, Stable-baselines3, ROS2
- **Optimization & Tools:** CUDA, NumPy, Pandas, OpenCV, SciPy
- **Machine Learning & AI:** Machine Learning, Multimodal Machine Learning, RAG, NLP, Finetuning, Distillation, LLM, Vision Language Models, Reinforcement Learning, Distributed Training,

EXPERIENCE

Fischell Department of Bioengineering - University of Maryland | Machine Learning Intern

Jun 2024 - Present

- Implemented YOLOv8 for dredger detection in a University of Maryland initiative to modernize oyster farming, later upgrading to RT-DETR which improved IoU by 6%.
- Developed human pose detection for walking style classification using MMPose and a CNN transformer, enhancing classification accuracy and contributing to improved gait analysis.

Sai Vamsi Industries | Machine Learning Engineer

Jan 2022 - May 2023

- Integrated a camera-based system using the YOLO framework to detect visual anomalies in press tool machines, showcasing applied computer vision and image understanding skills.
- Combined vibration data from Fluke 3561 FC Vibration Sensors with real-time image analysis to enhance defect detection by 5%, demonstrating practical software development and distributed training methodologies.

PROJECTS

Implementation of GPT and GPT 2 Models from Scratch

March 2025

- Developed a Generative Pretrained Transformer (GPT) model entirely from scratch using PyTorch, implementing core components such as tokenization, self-attention mechanisms, multi-head attention, positional embeddings, and transformer blocks.
- Reproduced a 124M parameter GPT-2 model, including architecture design, model training, fine-tuning, and inference. Optimized training pipelines, configured hyperparameters, and validated model performance against standard language benchmarks.

Text-to-Command Translation for Robot Navigation Using T5-small Transformer Model [\[GitHub Link\]](#)

Oct 2024 - Dec 2024

- Fine-tuned T5-Small to translate natural language commands into structured navigation plans with 98% test accuracy.
- Integrated Low-Rank Adaptation (LoRA) to reduce training parameters by 99.64% while retaining 92.5% test accuracy.
- Implemented the resulting NLP model with ROS2 and Gazebo to validate autonomous navigation on a TurtleBot3.

Transformers for 3D Object Detection in LiDAR Point Cloud [\[GitHub Link\]](#)

Nov 2024 - Dec 2024

- Developed and trained a custom transformer-based 3D object detection model for LiDAR point clouds on the KITTI dataset.
- Combined PointNet++ embeddings with specialized transformer encoders and introduced a novel loss function, achieving a 3D bounding box detection mAP IoU of 0.67—on par with leading benchmarks.

AI learns to Play MARIO Using Deep-Q Learning and SWIN Transformer [\[GitHub Link\]](#)

Mar 2024 - May 2024

- This project combines DQN with a SWIN Transformer to train Mario-playing AI agent.
- Enabling efficient interpretation of complex visuals and achieving a 500-moving average reward of 2700 in just 1250 episodes—far outperforming a model without SWIN.

Advanced Vision Systems for Autonomous Navigation Using YOLOv5 and Homography [\[GitHub Link\]](#)

Jan 2024 - May 2024

- Applied homography for dynamic route planning.
- Optical flow and YOLOv5 were used to detect obstacles and their velocities in real time, enabling adaptive navigation.
- By leveraging projective geometry to locate vanishing points and horizon lines, the robot achieved over 80% successful navigation.

Solved Airplane Passenger Problem Using RNN and LSTM models [\[GitHub Link\]](#)

Jan 2024 - Feb 2024

- Developed RNN and LSTM models for international airline passenger forecasting—leveraging advanced data preprocessing and time-series analysis—where the LSTM’s long-term dependency capture drove a 90% accuracy rate and markedly outperformed the RNN, highlighting its potential to optimize airline operations and resource management.